

Project

CNN for pneumonia categorizing analysis using shap



CNNs & Pytorch background

- No practical lectures on deep nns above 3-5 layers
- One short lab on pytorch in earlier course
- We will have “Neural networks” course in 1 year
- Course lectures: Basecase 1 layer CNN
 - No mention of channels, transfer learning, freezing or best practices
- Resnet18 : Nested convs, skip connections, batchnorms, pooling

So I built one of my own to understand our task



Sorry that was the wrong image

```
class hbblock(nn.Module):
    def __init__(self,in_channels, out_channels, identity_downsample=None, stride=1):
        #connections in, connections out, identity_downsample=we keep dimensions across skip connection stride=1 for convolution steps
        super(hbblock, self).__init__()
        self.expansion=4 #bottleneck block thing. Something about decreasing calculation size and then return to "normal" after a given block
        self.conv1= nn.Conv2d(in_channels, out_channels, kernel_size=1,stride=1, padding=0) #1st convolutional layer, doing normal jumps
        #kernel size= size of conv filter, 1 in this case is for cheap channel reduction
        self.bn1 = nn.BatchNorm2d(out_channels) #z scale output keeping output dimensions
        self.conv2= nn.Conv2d(out_channels, out_channels, kernel_size=3,stride=stride, padding=1) #double out channels to maintain size
        #2nd conv adds some padding and takes in stride argument in case end user wants to tune
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.conv3 = nn.Conv2d(out_channels, out_channels*self.expansion, kernel_size=1, stride=1, padding=0) #output will 4x the size after the bottleneck
        self.bn3 = nn.BatchNorm2d(out_channels*self.expansion)
        self.relu = nn.ReLU() #why did we learn of 10+ activation functions in last course when 90% is relu?
        self.identity_downsample = identity_downsample #changes x from skip connection to match block output
```

```
x+= identity #AND THIS IS WHERE SHAP DEEPEXPLAINER BROKE!
#It wants x = identity + x since that is not an in place modification but im not sure why that breaks shap
# reason for in place modification is so a new (BIG) tensor isnt created before the old one is deleted.
# So shap-ing a big nn will be extra intense
#but I think you can do in place while training and re define forward when its time for shap
x=self.relu(x)
return x
```

```
        self.in_channels=out_channels*4 #this will create a *4*4 (256) effect due to self.expansion in hbblock class
        # NOTE at this point I think one can visualize number of channels for a given a CNN with bottleneck layers sortof like a multimodal violin plot


        #building more bottleneck blocks
        for i in range(num_residual_blocks-1): #i since one block is already created above
            layers.append(block(self.in_channels, out_channels, identity_downsample=None, stride=1))
            #start of block: inchannels= 256, out=64; end of block: out=64*4=256
        return nn.Sequential(*layers) #unpack the list for pytorch

hbresnet50(img_channels=1, num_classes=2): #guide says img_channels=3 RGB and 1k classes but ill just start with ch=1,classes=2 for pneumoniamnist
return(hbResNet(hbblock, [3,4,6,3],img_channels, num_classes))

#The level of abstraction is actually insane here. Very cool
hbresnet101(img_channels=1, num_classes=2):
return(hbResNet(hbblock,[3,4,23,3],img_channels, num_classes))
```


First untuned run and CM images

- NUM_EPOCHS=5
 - Small since its pretrained
- lr=0.01
- Momentum=0.9
- BATCH_SIZE = 128
- Train/Test progress shows room for more epochs (assuming no frozen layers):
 - Epoch 1: Train 0.8736 Test 0.7324
 - Epoch 4: Train 0.9539 Test 0.8061
 - **Epoch 5: Train 0.9813 Test 0.8734**
- ROC curve show good results but room for adjusting threshold if we want more specificity (Optimal = TP-FP). This seems reasonable given the fact that a chest Xray implies high suspicion of pneumonia.
- CM confirms that we are indeed high up on the ROC curve.
 - More false positives than false negatives.




Shap first run

- Automatic “CM” index and image extraction lists
- Meaning we might get different images for shap analysis if the tuned NN classifies the images here.
 - Tuned shap image plots might differ wildly because of this
- TN looks very happy not having Pneumonia!
 - Or possibly doing a snatch




Gradient explainer

- In order: model_explainer 1, 0. Layer_1_explainer 1,0. Layer_4_explainer 1,0
- **Model pixels.** Detects the shapes around the lungs.
 - The false negative has one wrongly confident blob in the LH bottom of the image.
 - Big pixel spread considering decent accuracy for the first round
 - There also seem to be more spread on missclassified predictions
 - TN seem very confident across the whole NN
- **Layer 1.** You can see the patients right thorax! Probably fine to freeze this
- **Layer 4.** Highly interpolated due to 2x2 downsampling! Only color intensity relevant?



Bonus!

- In the sense that “I” sort of solved these 2 before project description changed from "Partition and deepexplainer" to "Gradient explainer"
- It is also impossible to discuss image classification with less than 1 cat picture
(<https://arxiv.org/abs/2012.01768>)



Partition explainer


- Grabs one image. Blurs a given pixel or superpixel, see what model does.
Repeat until we get "marginal (super)pixel contribution" == Img shap values.
Does not need a "shap training set" the way one needs for a normal regression model for example
Only does local image analysis

- Params:
`batch_size=100`
`n_evals = 300`
`masker = "blur 28 28", img.shape`

- Images:
TP FP
FN TN

- TN also shows more confidence
confirming gradientexplainers results
even if partition is a local explainer vs
Gns global explanation


- Surprisingly both explainers are not very confident in
TP predictions even if ROC is fairly high compared to
optima



Deepexplainer

- Turns out deepexplaner has had known issues with pytorch since january 2024.
It took me a bunch of digging and a ton of vibecoding before "I" got a working solution.
The main issue was that during backprop, torch does in-place modification of gradients in tensor. Along with in place modification of relu.
We never had a lecture on tensors. So from what i read up on the topic its more or less NumPy arrays with built in gradients
I never understood the issue until I built the "Big resnet from scratch"
But it took several hours to vibe-solve and even more to understand.
- Depexplainer is a slightly less handwavy nutshell than the 10+ articles I read:
Approximates shapvals by measuring activation of each individual neuron
Each matrix or normal layer neuron is treated as a single feature

- Images: pred 1, 0
TP FP
FN TN




The method name is misleading

- Deepexplainer uses neurons and conv/pool indices to **approximate shap vals**
- It works by comparing each neuron's activation for your image to its activation for a "neutral" baseline image.
The difference is then calculated as weight * (change from neutral) back through the network's connections to the input pixels.
- Inference: Deepexplainer should get worse as nn gets deeper
- This is almost as bad as skip connections!
-Functionally they boost the gradient by 1



Freezing and tuning

- No prior explanation on this topic so this is mostly inference and guesswork.
Main ref is stanford cs231 article "4 scenarios"
- ***Small dataset – similar to original: Freeze all conv and train the final layer. Huge overfitting risk.**
- Large dataset – similar to original: Can probably work through the whole NN. Be empirical and mindful of overfitting.
- ***Small dataset – dissimilar to original: Either just train final layer, train middle layers (maintain edge detection filters in the beginning) or both.**
-This falls in line with gradient explainers layer 4 plot. Its only "certain" on TN. The other shaps are not confident.
- Big dataset – dissimilar to the original: If you have money and time: Do a new nn from scratch. Save money by using a dissimilar pretrained nn anyway.
- LR and momentum should generally be smaller than baseline since weights might be too perturbed otherwise.
-Analogy in my head: We are running the final few epochs and lr_decay=True
- Couldn't find anything on batch size that is on a level for a 2nd year undergrad to understand.
Inferences:
 - Small dataset and frozen layers means we only adjust a few weights and increase the risk of overfitting.
 - Use smaller learning rate to counter overfitting. Increased risk of stuck in local minima.
 - Counter local minima by using smaller batches for stochastic GD
- Adjusting epochs last and go right to the overfitting edge
- Will mostly look at acc scores while tuning






Freezing diary

- Starting hyperparameters: Batch_size = 128/2, Lr= 0.01/3, Momentum =0.9/3, Epochs=5
- Original results for reference: Train 0.9813 Test 0.8734
- fc unfreeze: train 0.8734, test 0.7837
 - Too conservative but smaller overtraining gap.
- fc, l4 unfreeze: train 0.9751, test 0.8333
 - Too aggressive Will cut lr and momentum by half and keep freezing
- fc, l4 unfreeze lr=0.01/3/2, momentum=0.9/3/2: train 0.9431 test 0.8237
 - Much better on overfitting but same diff as updating the whole network
- fc, l4, l3: train 0.9743, test 0.8333
 - Back to overfitting will do another cut on lr/momentum but at this stage I want to go back to just freezing l4 and fc
- fc, l4, l3 lr=0.01/3/2/2, momentum=0.9/3/2/2: Train 0.9480 test 0.8285
 - The relative overfitting is consistent. Will backtrack to layer 4 and fc with lower lr along with decreasing batch size and increasing epochs
- fc, l4 batch size=128/2/2batch epochs=10: Train 0.9446, test 0.8301
 - Several iterations of fc, l4 unfreezes never coming close to original full nn output.
- **Smallest distance between train and test was simply fc unfrozen.**
 - Rerunning with lower lr=0.0016 and higher epochs =15.
- Documentation will be less detailed due to marginal changes in absolute values and always getting stuck in local minima.
 - Train and test both bounced back and forth no matter what tuning combination used.
- Results seem to align with the small dataset leading to overfitting risk mentioned in CS231 freezing article.
- No experiment was better than the default values mentioned at the start.
- Tuning experiments on the whole network
- Batch size= 64, lr=0.05, momentum=0.045, epochs=8 ended up with 0.9964 test 0.8750.
 - A small gain and a very small increase of the overfitting gap.
- **This is probably a dataset size limitation.**
- **A possible solution would be to perturb images but time budget has already been exceeded by a large magnitude.**




Changes of outputs from tuning?

- Batch size= 64, lr=0.05, momentum=0.045, epochs=8 ended up with train 0.9964 test 0.8750.
- No change in CM index images
-TN is still happy!
- No statistical difference! Chi² p-value=1!
- Tuning increased accuracy but changed the sensitivity/specificity tradeoff in a clinically unfavorable direction due to decreased sensitivity
- This is reflected in the ROC “optimal” which is marginally lower than before ($0.988 \rightarrow 0.958$)



Gradient explainer

- Hard to quantify images compared to clf or regression shap values but eyeballing implies
- **Model confusion!**
 - Confidence increased for True negative in both directions, indicating confusion.
 - The other CM colors are weaker, showing a decrease in confidence
- **Layer 1** has shifted shapes slightly
 - But look overall less confident. Local minima explanation!?
- **Layer 4** though has done some major updating!
 - Pushing FP image towards going to the correct TN classification.
 - But at the cost of getting less confident in False negative (lost some pink in the bottom)




Partition explainer

- Images:


TP FP

FN TN



- No change in params for partition explainer

- TP:** More confident
-good!



- FP:** Less confident
-also good!


- TN:** Less confident
-bad

- FN:** Looks like its using the LH clavicle pixels more than lungs
-real bad!
-Tuning seemingly changed too much of the edge detection
-And probably explains why the model lost sensitivity

- Possible solution:** Gradual freezing based on these results.
-Starting with 1 or 2 epochs on all layers and then freezing one layer every 2 epochs to avoid changing feature detection too much
-Way outside skill and time constraints



Deepexplainer

- Images: pred 1, 0
TP FP
FN TN
- **TP:** Seems overall better at following the lung outline
- **FP:** Slightly more blue on RH lung field so moving in the correct direction
- **TN:** Slightly less confident?
 - Hard to quantify red pixels
- **FN:** Slightly less confident
 - Pushing towards the correct direction.
- Overall follows partition explainer
-but visualized differently



Final thoughts on shapvals

- Overall mixed and tiny, nuanced results.
 - Not surprising given lack of significant difference pre and post tuning
 - Stumbled on too good starting hyperparams!
- Doing gradient explainer with several layers seem best for NN freeze tuning
 - Loads of work though
- Really dislike not having quantification the same way .force_plot does
 - "Err I think this is a deeper red?"
- Also: Color blindness?!



Challenges working with Shap

- <1/2 of image classification sample code is on keras.
Very different from torch requirements since it more or less outputs correct shapes by default from what I can deduce.
- Warning messages in both “getting_started” and shap sample notebooks indicate badly maintained example code:
 - UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
 - UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future.
- “Resnet50 partitionexplainer”
 - No warning signs on example code for runtime or memory usage
 - Lack of code comments indicates users know what they are doing
 - Need ~32gb gpu minimum
 - Or risk runtime disconnecting on google collab
 - OOM errors produce TensorFlow stacktraces (remember, we haven't learned about tensors yet)
- “Explain an Intermediate Layer of VGG16 on ImageNet (PyTorch)”
 - Output is...Not optimal for learning the library
- Deepexplainer error even asked me to post the code on github:

AssertionError: The SHAP explanations do not sum up to the model's output! This is either because of a rounding error or because an operator in your computation graph was not fully supported. If the sum difference of %f is significant compared to the scale of your model outputs, please post as a [github issue](#), with a reproducible example so we can debug it. Used framework: [pytorch](#) - Max. diff: 0.7187470377143654 - Tolerance: 0.01

